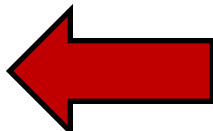


3D Viewing, Orthographic and Perspective Projection

COMP557

Paul Kry

Roadmap


- Transformations (FCA 6) A1
- Scene graphs (FCA 12.2) A2
- Viewing and Projection (FCA 7)  A3
- Meshes, simplification (FCA 12.1 + pdf)
- Subdivision (pdf + notes) A4
- ...

May insert material on graphics pipeline / rasterization, clipping /culling (FCA 8), light and shadow (FCA 4.5, 10), before meshes.

Recall: *Image order and Object order*

- Image order: “backward” approach
 - start from pixel
 - ask what part of scene projects to pixel
 - explicitly construct the ray corresponding to the pixel
- Object order: “forward” approach
 - start from a point in 3D
 - compute its projection into the image
- matrix transformations critical for object order approach:
 - combines seamlessly with coordinate transformations used to position camera and model
 - ultimate goal: single matrix operation to map any 3D point to its correct screen location.


Image Order Approach to Viewing



[Durand]

- Ray generation produces rays, not points in scene
- Cast a ray at every pixel and find points of intersection
- This is called *ray tracing*

Object Order Approach to Viewing



Projection (left) and rasterization (right) of a triangle.


- Inverting the ray tracing process requires division for the perspective case
- Once triangle vertices are known in screen coordinates the triangle can be filled in (rasterization... more on this later)

Mathematics of projection

- Always work in eye coords
 - assume eye point at **0** and plane perpendicular to z
- Orthographic case
 - a simple projection: just toss out z
- Perspective case: scale diminishes with z
 - and increases with d


Pipeline of transformations

- Standard sequence of transforms



Camera / Eye Coordinates

- We have discussed object to world transformations
- Need world to camera transformation
 - In viewing, we typically know:
 - Where the camera is
 - What we want to look at
 - Which way is up
 - Need a rigid transformation (rotation and translation)
 - How many degrees of freedom?
 - How to compute it?
 - Can use gluLookat



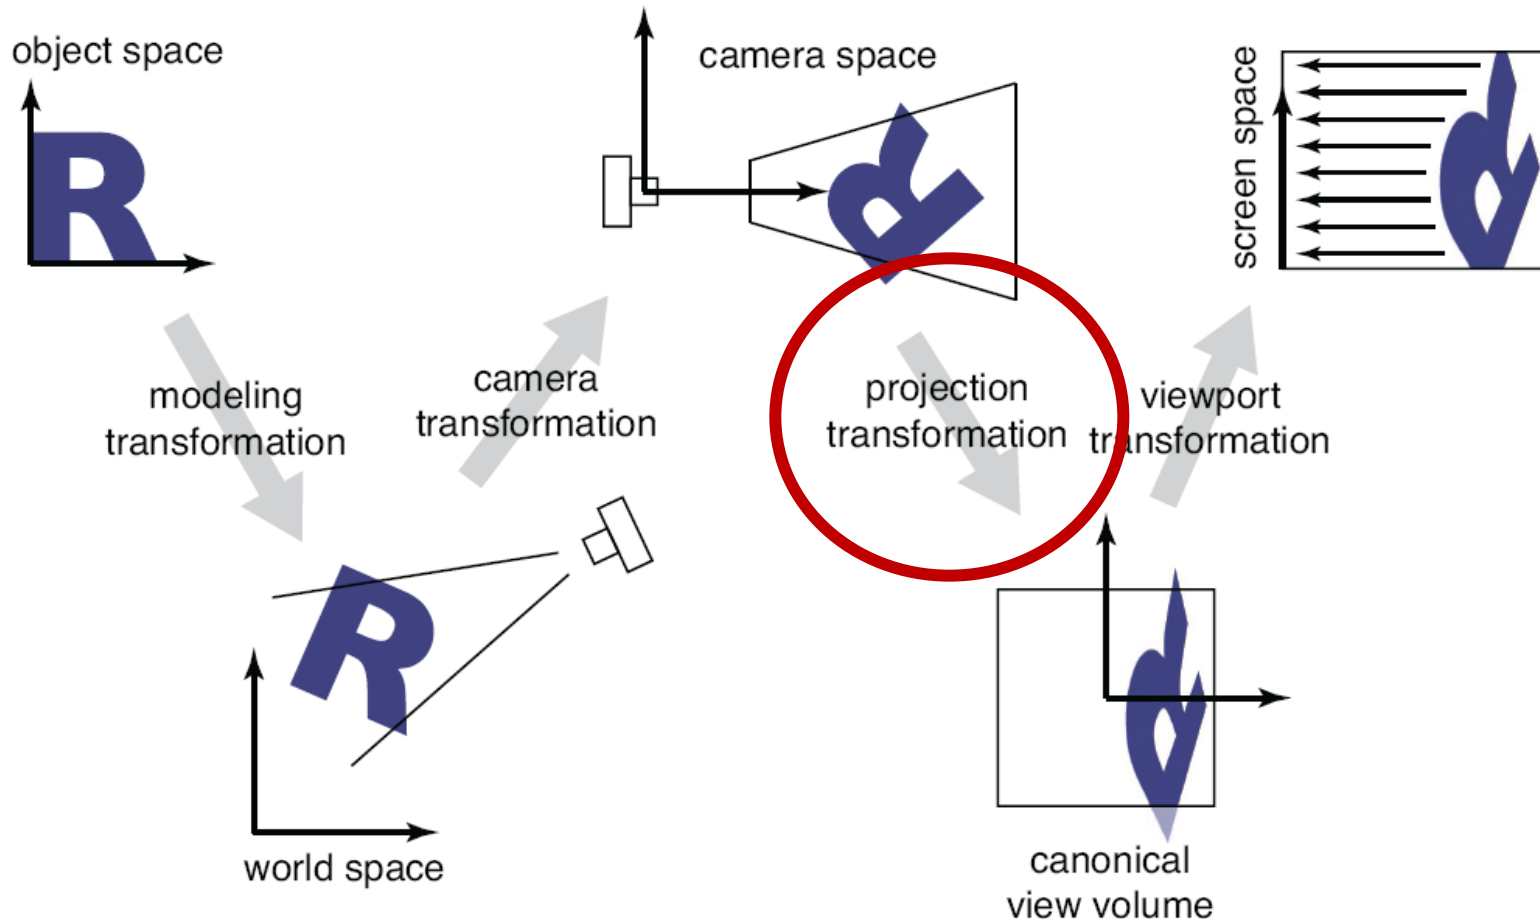
[Todo]

“Lookat” Transformation

- uvn or uvw commonly used for xyz camera axes
- Compute transform from points e , l , and vector V_{up}
 - e is the eye point, or view reference point (vrp)
 - l is the lookat point
- Separate the rotation R and translation T
 - What order to compose? What is easiest?
 - What is T ?
 - What is R ?
- We want $\begin{bmatrix} \mathbf{u} & \mathbf{v} & \mathbf{w} & \mathbf{e} \\ 0 & 0 & 0 & 1 \end{bmatrix}^{-1}$

Pipeline of transformations


- Standard sequence of transforms



Projection


- How to form an image by ***planar perspective projection***?

In computer graphics,
projection plane is in front
of center of projection



[Source unknown]

Parallel projection: orthographic




But lets start with something simpler... orthographic projection to implement orthographic, just toss out z .


$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} x \\ y \\ 1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Pipeline of transformations

- Standard sequence of transforms



View volume: orthographic




Viewing a cube of size 2

- Start by looking at a restricted case: the *canonical view volume*
- It is the cube $[-1,1]^3$, viewed from the z direction
- Matrix to project it into a square image in $[-1,1]^2$ is trivial:

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Viewing a cube of size 2


- To draw in image, need coordinates in pixel units
- Suppose n_x pixels in x direction and n_y pixels in y direction



Pixel size in canonical view volume is $2/n_x$ by $2/n_y$
Center is at $\frac{1}{2}$ pixel width and height away from $(-1, -1)$
i.e., $(-1+1/n_x, -1+1/n_y)$ maps to $(0,0)$ integer pixel location
 $(1-1/n_x, 1-1/n_y)$ maps to (n_x-1, n_y-1) integer pixel location

Windowing transforms

- This transformation is worth generalizing
 - take one axis-aligned rectangle or box to another
 - a useful transformation chain




$$\begin{bmatrix} 1 & 0 & x'_l \\ 0 & 1 & y'_l \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \frac{x'_h - x'_l}{x_h - x_l} & 0 & 0 \\ 0 & \frac{y'_h - y'_l}{y_h - y_l} & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & -x_l \\ 0 & 1 & -y_l \\ 0 & 0 & 1 \end{bmatrix}$$

$$\begin{bmatrix} \frac{x'_h - x'_l}{x_h - x_l} & 0 & \frac{x'_l x_h - x'_h x_l}{x_h - x_l} \\ 0 & \frac{y'_h - y'_l}{y_h - y_l} & \frac{y'_l y_h - y'_h y_l}{y_h - y_l} \\ 0 & 0 & 1 \end{bmatrix}$$

[Shirley3e f. 6-16; eq. 6-6]

Viewport transformation



$$\begin{bmatrix} x_{\text{screen}} \\ y_{\text{screen}} \\ 1 \end{bmatrix} = \begin{bmatrix} \frac{n_x}{2} & 0 & \frac{n_x - 1}{2} \\ 0 & \frac{n_y}{2} & \frac{n_y - 1}{2} \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} x_{\text{canonical}} \\ y_{\text{canonical}} \\ 1 \end{bmatrix}$$

$(-1+1/n_x, -1+1/n_y)$ maps to $(0,0)$

$(1-1/n_x, 1-1/n_y)$ maps to (n_x-1, n_y-1)


Viewport transformation

- In 3D, carry along z for the ride
 - one extra row and column

$$\mathbf{M}_{vp} = \begin{bmatrix} \frac{n_x}{2} & 0 & 0 & \frac{n_x-1}{2} \\ 0 & \frac{n_y}{2} & 0 & \frac{n_y-1}{2} \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Orthographic projection

- First generalization: different view rectangle
 - retain the minus-z view direction



- specify view by left, right, top, bottom (as in RT)
- also near, far

Clipping planes

- Recall...
 - Object-order rendering considers each object in turn
 - i.e., forward rendering, rasterization
 - Image-order rendering considers each pixel in turn
 - i.e., backward rendering, ray tracing
- In object-order systems we **always** use at least two *clipping planes* that further constrain the view volume
 - near plane: parallel to view plane; things between it and the viewpoint will not be rendered
 - far plane: also parallel; things behind it will not be rendered
- These planes are:
 - partly to *remove unnecessary stuff* (e.g., behind the camera)
 - but really to *constrain the range of depths* (we'll see why later)

Orthographic projection

- We can implement this by mapping the view volume to the canonical view volume.
- This is just a 3D windowing transformation!


$$\mathbf{M}_{\text{orth}} = \begin{bmatrix} \frac{x'_h - x'_l}{x_h - x_l} & 0 & 0 & \frac{x'_l x_h - x'_h x_l}{x_h - x_l} \\ 0 & \frac{y'_h - y'_l}{y_h - y_l} & 0 & \frac{y'_l y_h - y'_h y_l}{y_h - y_l} \\ 0 & 0 & \frac{z'_h - z'_l}{z_h - z_l} & \frac{z'_l z_h - z'_h z_l}{z_h - z_l} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

$$\mathbf{M}_{\text{orth}} = \begin{bmatrix} \frac{2}{r-l} & 0 & 0 & -\frac{r+l}{r-l} \\ 0 & \frac{2}{t-b} & 0 & -\frac{t+b}{t-b} \\ 0 & 0 & \frac{2}{n-f} & -\frac{n+f}{n-f} \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

Camera and modeling matrices

- We worked out all the preceding transforms starting from eye coordinates
 - before we do any of this we need to transform into that space
- Transform from world (canonical) to eye space is traditionally called the *viewing matrix*
 - Easy for us to compute the matrix F_c which takes us from eye space to canonical space, but here we use the inverse F_c^{-1}
- Remember that geometry would originally have been in the object's local coordinates; transform into world coordinates is called the *modeling matrix*, M_m
- Note some systems (e.g., OpenGL) combine the two into a *modelview* matrix and just skip world coordinates

Viewing transformation



the camera matrix rewrites all coordinates in eye space


Orthographic transformation chain

- Start with coordinates in object's local coordinates
- Transform into world coords (modeling transform, M_m)
- Transform into eye coords (camera xf., $M_{\text{cam}} = F_c^{-1}$)
- Orthographic projection, M_{orth}
- Viewport transform, M_{vp}

$$\mathbf{p}_s = \mathbf{M}_{\text{vp}} \mathbf{M}_{\text{orth}} \mathbf{M}_{\text{cam}} \mathbf{M}_m \mathbf{p}_o$$

$$\begin{bmatrix} x_s \\ y_s \\ z_c \\ 1 \end{bmatrix} = \begin{bmatrix} \frac{n_x}{2} & 0 & 0 & \frac{n_x-1}{2} \\ 0 & \frac{n_y}{2} & 0 & \frac{n_y-1}{2} \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \frac{2}{r-l} & 0 & 0 & -\frac{r+l}{r-l} \\ 0 & \frac{2}{t-b} & 0 & -\frac{t+b}{t-b} \\ 0 & 0 & \frac{2}{n-f} & -\frac{n+f}{n-f} \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \mathbf{u} & \mathbf{v} & \mathbf{w} & \mathbf{e} \\ 0 & 0 & 0 & 1 \end{bmatrix}^{-1} \mathbf{M}_m \begin{bmatrix} x_o \\ y_o \\ z_o \\ 1 \end{bmatrix}$$

Planar Perspective projection



With the Projection plane at distance d ,
note the similar triangles:

$$\frac{y'}{d} = \frac{y}{-z}$$

$$y' = -dy/z$$

- y' is the foreshortened version of y , that is, it is smaller by a factor $-d/z$
- We can think of dividing by negative z so that y' has the same sign as y !

Homogeneous coordinates revisited

- Perspective requires division
 - that is not part of affine transformations
 - in affine, parallel lines stay parallel
 - therefore no vanishing point
 - therefore no rays converging on viewpoint
- “True” purpose of homogeneous coords: projection

Homogeneous coordinates revisited

- Introduced $w = 1$ coordinate as a placeholder

$$\begin{bmatrix} x \\ y \\ z \end{bmatrix} \rightarrow \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

– used as a convenience for unifying translation with linear

- Can also allow arbitrary w


$$\begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} \sim \begin{bmatrix} wx \\ wy \\ wz \\ w \end{bmatrix}$$

Implications of w

$$\begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} \sim \begin{bmatrix} wx \\ wy \\ wz \\ w \end{bmatrix}$$

- All scalar multiples of a 4-vector are equivalent
- When w is not zero, can divide by w
 - therefore these points represent “normal” affine points
- When w is zero, it's a point at infinity, i.e., a direction
 - can think of this as the point where parallel lines intersect
 - can also think of it as the vanishing point


Perspective projection




to implement perspective, just move z to w:

$$\begin{bmatrix} x' \\ y' \\ 1 \end{bmatrix} = \begin{bmatrix} -dx/z \\ -dy/z \\ 1 \end{bmatrix} \sim \begin{bmatrix} dx \\ dy \\ -z \end{bmatrix} = \begin{bmatrix} d & 0 & 0 & 0 \\ 0 & d & 0 & 0 \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

View volume: perspective



View volume: perspective (clipped)



Carrying depth through perspective

- Perspective has a varying denominator—can't preserve depth!
- Compromise: preserve order, and depth on near and far planes

$$\begin{bmatrix} x' \\ y' \\ z' \\ 1 \end{bmatrix} \sim \begin{bmatrix} \tilde{x} \\ \tilde{y} \\ \tilde{z} \\ -z \end{bmatrix} = \begin{bmatrix} d & 0 & 0 & 0 \\ 0 & d & 0 & 0 \\ 0 & 0 & a & b \\ 0 & 0 & -1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

– that is, choose a and b so that $z'(n) = n$ and $z'(f) = f$.

$$\tilde{z}(z) = az + b$$

$$z'(z) = \frac{\tilde{z}}{-z} = \frac{az + b}{-z}$$

want $z'(n) = n$ and $z'(f) = f$

result: $a = -(n + f)$ and $b = nf$ (try it)

Official perspective matrix

- Use near plane distance as the projection distance
- Let n be **negative** and $d = -n$
- Scale by -1 to have fewer minus signs
 - The matrix is different but this does not change the projective transformation

$$\mathbf{P} = \begin{bmatrix} n & 0 & 0 & 0 \\ 0 & n & 0 & 0 \\ 0 & 0 & n + f & -fn \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

Potential Confusion:
In OpenGL the near and far plane are specified as positive numbers giving the distance along the negative z axis!

Questions

- Questions to better understand the action of a 4x4 homogenous planar perspective transformation matrix
 - What happens to lines through the origin ?
 - They become parallel
 - What happens to vectors $(x,y,z,0)$?
 - They map to plane $z=\text{near}+\text{far}$
 - What happens to points on ***plane with z normal*** at center of projection ?
 - Map to points at infinity (all vectors come from the homogeneous representations of points on this plane at center of projection)
 - What happens to points at $z = (\text{near}+\text{far})/2$?
 - $(n*\text{near}+f*\text{far}) / (n+f)$, that is, z is not preserved !! However, order is! (has implications for z depth precision)
 - What happens to points behind the camera ?
 - It goes in front of the camera



Questions (the easier version)


- Given the following projection matrix, with near at 1 and far at 3, answer the following by interpreting the result in non-homogeneous coordinates

$$P = \begin{pmatrix} -1 & 0 & 0 & 0 \\ 0 & -1 & 0 & 0 \\ 0 & 0 & -4 & -3 \\ 0 & 0 & 1 & 0 \end{pmatrix}$$



- What happens to points on the near plane?
- What happens to points on the far plane?
- What happens to points half way between?
- What happens to points behind the camera?

Aside: What goes where?




camera coordinates

$$z = f_0 + f_1$$

$$z = f_1$$

$$z = f_0$$

$$z = 0$$



(non-normalized) projection coordinate

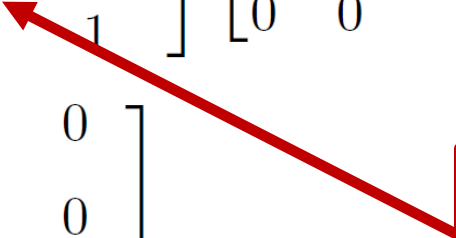
This has implications for clipping (i.e., discarding) geometry!

Perspective projection matrix

- Need perspective projection to produce points in the canonical view volume, so combine with an orthographic projection matrix (windowing transform)

$$\mathbf{M}_{\text{per}} = \mathbf{M}_{\text{orth}} \mathbf{P}$$

$$= \begin{bmatrix} \frac{2}{r-l} & 0 & 0 & -\frac{r+l}{r-l} \\ 0 & \frac{2}{t-b} & 0 & -\frac{t+b}{t-b} \\ 0 & 0 & \frac{2}{n-f} & -\frac{n+f}{n-f} \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} n & 0 & 0 & 0 \\ 0 & n & 0 & 0 \\ 0 & 0 & n+f & -fn \\ 0 & 0 & 1 & 0 \end{bmatrix}$$

$$= \begin{bmatrix} \frac{2n}{r-l} & 0 & \frac{l+r}{l-r} & 0 \\ 0 & \frac{2n}{t-b} & \frac{b+t}{b-t} & 0 \\ 0 & 0 & \frac{f+n}{n-f} & \frac{2fn}{f-n} \\ 0 & 0 & 1 & 0 \end{bmatrix}$$


Is n-f a bug?

Is it not f-n ??

The convention is that -near maps to -1 and -far maps to 1.


Perspective transformation chain

- Transform into world coords (modeling transform, M_m)
- Transform into eye coords (camera xf., $M_{\text{cam}} = F_c^{-1}$)
- Perspective matrix, P
- Orthographic projection, M_{orth}
- Viewport transform, M_{vp}

$$\mathbf{p}_s = \mathbf{M}_{\text{vp}} \mathbf{M}_{\text{orth}} \mathbf{P} \mathbf{M}_{\text{cam}} \mathbf{M}_m \mathbf{p}_o$$


$$\begin{bmatrix} x_s \\ y_s \\ z_c \\ 1 \end{bmatrix} = \begin{bmatrix} \frac{n_x}{2} & 0 & 0 & \frac{n_x-1}{2} \\ 0 & \frac{n_y}{2} & 0 & \frac{n_y-1}{2} \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} \frac{2}{r-l} & 0 & 0 & -\frac{r+l}{r-l} \\ 0 & \frac{2}{t-b} & 0 & -\frac{t+b}{t-b} \\ 0 & 0 & \frac{2}{n-f} & -\frac{n+f}{n-f} \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} n & 0 & 0 & 0 \\ 0 & n & 0 & 0 \\ 0 & 0 & n+f & -fn \\ 0 & 0 & 1 & 0 \end{bmatrix} \mathbf{M}_{\text{cam}} \mathbf{M}_m \begin{bmatrix} x_o \\ y_o \\ z_o \\ 1 \end{bmatrix}$$

OpenGL view frustum: orthographic



Note OpenGL puts the near and far planes at $-n$ and $-f$ so the user should give positive numbers

OpenGL view frustum: perspective



Note OpenGL puts the near and far planes at $-n$ and $-f$ so that the user should give positive numbers


`gluPerspective(fovy, aspect, near, far)`


`glFrustum(left, right, bottom, top, near, far)`


$\underbrace{\hspace{10em}}$
Defined on near plane


Frustum applications

- Shifted perspective
- Tiled rendering (e.g., render very high resolutions)
- 3D viewing (i.e., left eye right eye)
- Depth of field (i.e., accumulating multiple render passes)









Review and more information

- Textbook chapter 7
 - Viewing and projection