

Ray Tracing II

**These slides are not so exciting,
we skiped over most**


Topics

- Transformations in ray tracing
 - Transforming objects
 - Transformation hierarchies
- Ray tracing acceleration structures
 - Bounding volumes
 - Bounding volume hierarchies
 - Uniform spatial subdivision
 - Adaptive spatial subdivision

Transforming objects

- In modeling, we've seen the usefulness of transformations
 - How to do the same in RT?
- Take spheres as an example: want to support transformed spheres
 - Need a new Surface subclass
- Option 1: transform sphere into world coordinates
 - Write code to intersect arbitrary ellipsoids
- Option 2: transform ray into sphere's coordinates
 - Then just use existing sphere intersection routine

Intersecting transformed objects



[Shirley 2000]

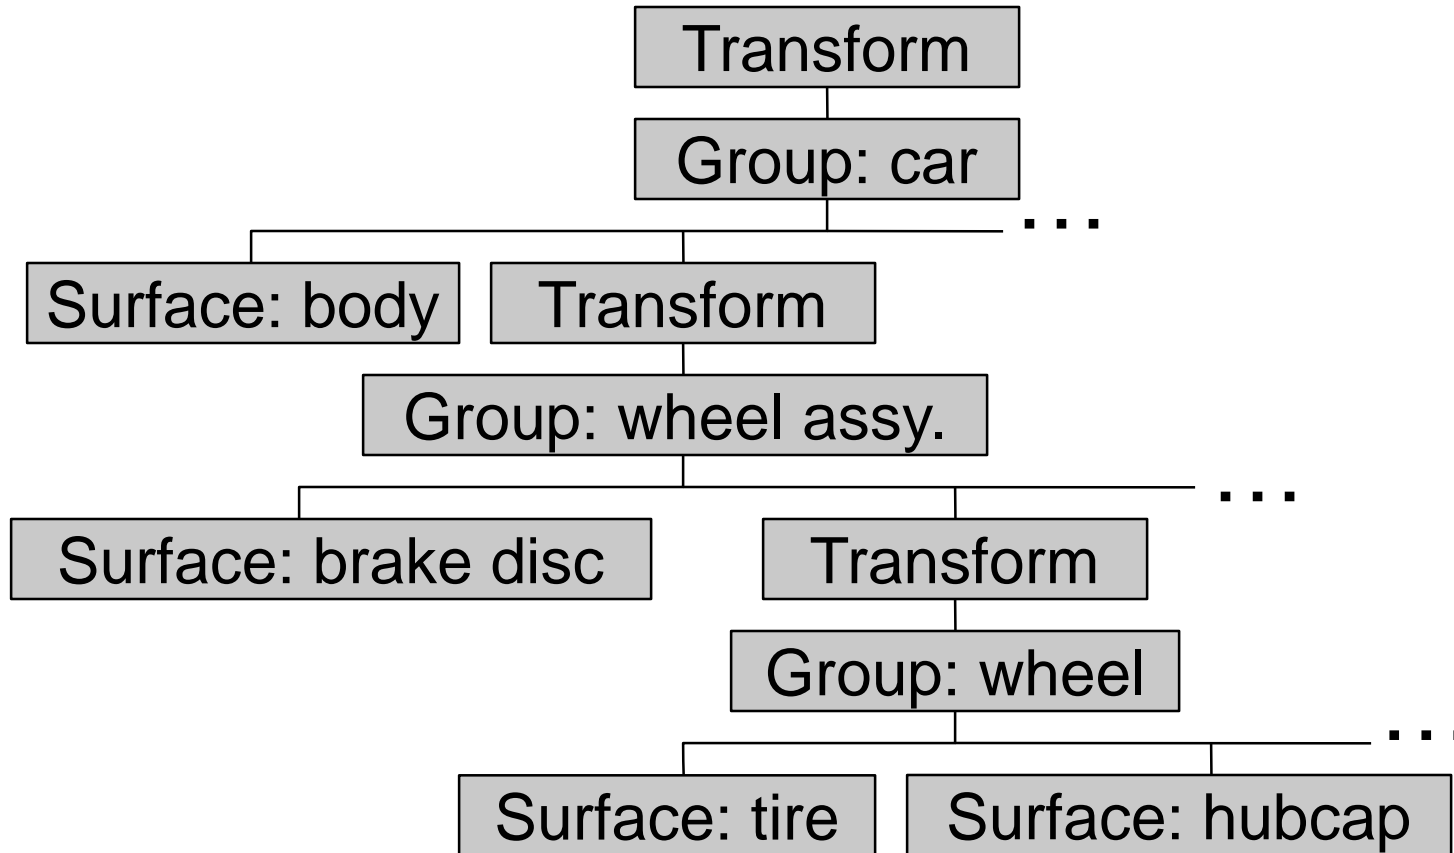
Implementing RT transforms

- Create wrapper object “TrasformedSurface”
 - Has a transform T and a reference to a surface S
 - To intersect:
 - Transform ray to local coords (by inverse of T)
 - Call `surface.intersect`
 - Transform hit data back to global coords (by T)
 - Intersection point
 - Surface normal
 - Any other relevant data (maybe none)

Groups, transforms, hierarchies

- Often it's useful to transform several objects at once
 - Add “SurfaceGroup” as a subclass of Surface
 - Has a list of surfaces
 - Returns closest intersection
 - Opportunity to move ray intersection code here to avoid duplication
- With TransformedSurface and SurfaceGroup you can put transforms below transforms
 - This gives us a transformation hierarchy.

A transformation hierarchy

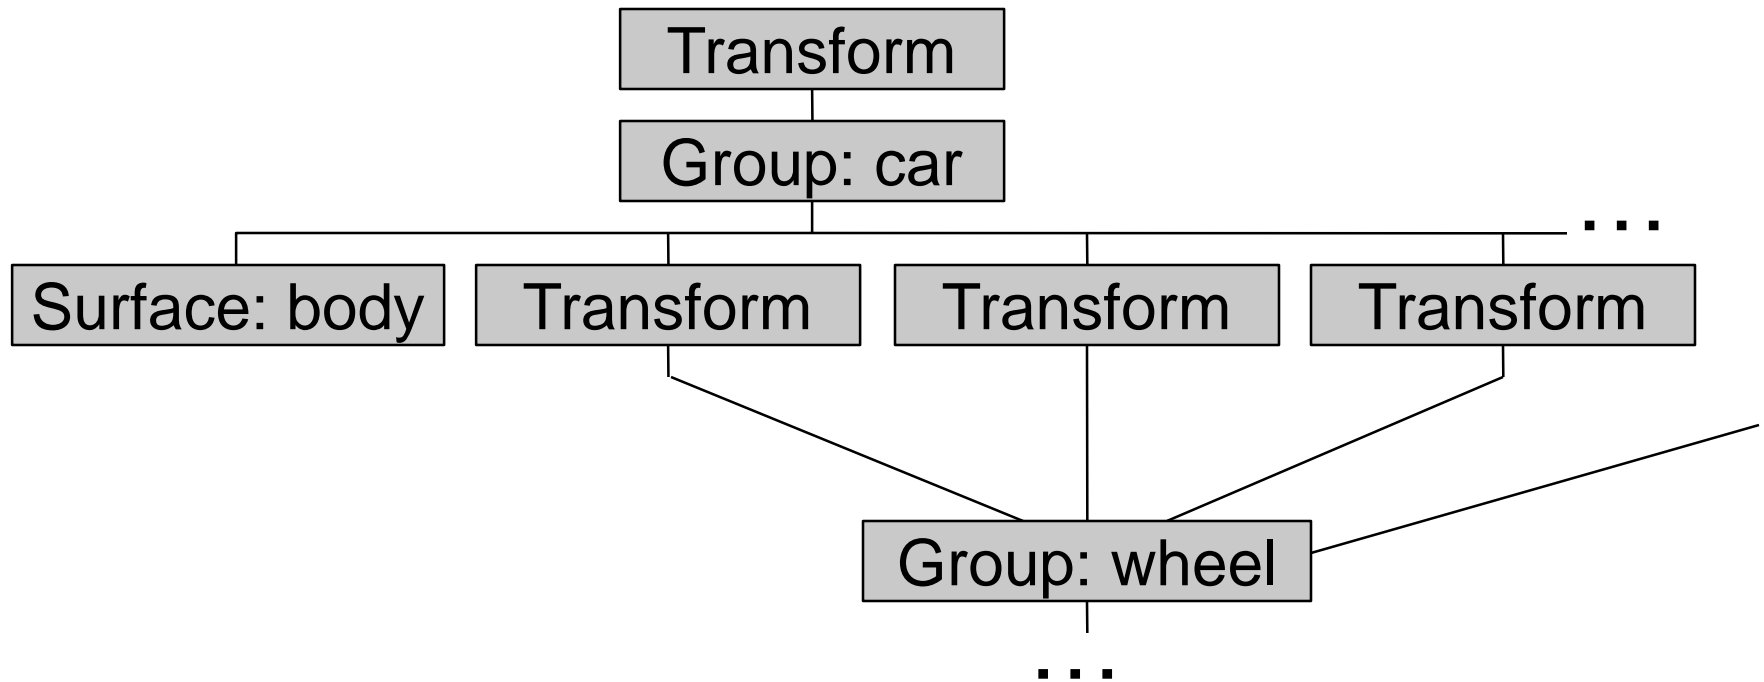


- Common optimization: merge transforms with groups
 - This is how we did it in the modeler assignment

Instancing

- Anything worth doing is worth doing n times
- If we can transform objects, why not transform them several ways?
 - Many models have repeated subassemblies
 - Mechanical parts (wheels of car)
 - Multiple objects (chairs in classroom, ...)
 - Nothing stops you from creating two TransformedSurface objects that reference the same Surface
 - Allowing this makes the transformation tree into a DAG
 - (directed acyclic graph)
 - Mostly this is transparent to the renderer

Hierarchy with instancing



Hierarchies and performance


- Transforming rays is expensive
 - minimize tree depth: flatten on input
 - push all transformations toward leaves
 - triangle meshes may do best to stay as group
 - transform ray once, intersect with mesh
 - internal group nodes still required for instancing
 - can't push two transforms down to same child!

Ray tracing acceleration

- Ray tracing is slow. This is bad!
 - Ray tracers spend most of their time in ray-surface intersection methods
- Ways to improve speed
 - Make intersection methods more efficient
 - Yes, good idea. But only gets you so far
 - Call intersection methods fewer times
 - Intersecting every ray with every object is wasteful
 - Basic strategy: efficiently find big chunks of geometry that definitely do not intersect a ray

Bounding volumes

- Quick way to avoid intersections: bound object with a simple volume
 - Object is fully contained in the volume
 - If it doesn't hit the volume, it doesn't hit the object
 - So test bvol first, then test object if it hits



Bounding volumes

- Cost: slightly more for hits and near misses, much less for far misses
- Worth doing? It depends:
 - Cost of bvol intersection test should be small
 - Therefore use simple shapes (spheres, boxes, ...)
 - Cost of object intersect test should be large
 - Bvols most useful for complex objects
 - Tightness of fit should be good
 - Loose fit leads to extra object intersections
 - Tradeoff between tightness and bvol intersection cost

Implementing bounding volume

- Just add new Surface subclass, “BoundedSurface”
 - Contains a bounding volume and a reference to a surface
 - Intersection method:
 - Intersect with bvol, return false for miss
 - Return surface.intersect(ray)
 - Like transformations, common to merge with group
 - This change is transparent to the renderer (only it might run faster)
- Note that all Surfaces will need to be able to supply bounding volumes for themselves


If it's worth doing, it's worth doing hierarchically!

- Bvols around objects may help
- Bvols around groups of objects will help
- Bvols around parts of complex objects will help
- Leads to the idea of using bounding volumes all the way from the whole scene down to groups of a few objects
- Meshes are a special case
 - We don't add individual triangles into the DAG
 - We probably want to do something hierarchically at the level of an individual triangle soup!


Implementing a bvol hierarchy

- A BoundedSurface can contain a list of Surfaces
- Some of those Surfaces might be more BoundedSurfaces
- Voilà! A bounding volume hierarchy
 - And it's all still transparent to the renderer

BVH construction example



BVH ray-tracing example




Choice of bounding volumes

- Spheres
 - easy to intersect, not always so tight,
 - Tricky to compute tight bounds (google Welzl's smallest enclosing disk algorithm)
- Axis-aligned bounding boxes (AABBs) –
 - easy to intersect, often tighter (esp. for axis-aligned models)
- Oriented bounding boxes (OBBs) –
 - easy to intersect (but cost of transformation), tighter for arbitrary objects
- Computing the bvols
 - For primitives -- generally pretty easy
 - For groups -- not so easy for OBBs (to do well)
 - For transformed surfaces -- not so easy for spheres

Axis aligned bounding boxes

- Probably easiest to implement
- Computing for primitives
 - Cube: very straightforward!
 - Sphere, cylinder, etc.: pretty obvious
 - Groups or meshes: min/max of component parts
- AABBs for transformed surface
 - Easy to do conservatively: bbox of the 8 corners of the bbox of the untransformed surface
- How to intersect them
 - Treat them as an intersection of slabs (see Shirley)

Intersecting boxes



$$t \in [t_{xmin}, t_{xmax}]$$

$$t \in [t_{ymin}, t_{ymax}]$$

$$t \in [t_{xmin}, t_{xmax}] \cap [t_{ymin}, t_{ymax}]$$

Building a hierarchy


- Usually do it top-down
- Make bbox for whole scene, then split into (maybe 2) parts
 - Recurse on parts
 - Stop when there are just a few objects in your box

Building a hierarchy

- How to partition?
 - Ideal: clusters
 - Practical: partition along axis
 - Median partition
 - More expensive
 - More balanced tree
 - Center partition
 - Less expensive, simpler
 - Unbalanced tree, but that may actually be better


Regular space subdivision

- An entirely different approach: uniform grid of cells




Regular grid example

- Grid divides space, not objects




Traversing a regular grid



Non-regular space subdivision

- *k*-d Tree
 - subdivides space, like grid
 - adaptive, like BVH



Implementing acceleration structures


- Conceptually simple to build acceleration structure into scene structure
- Better engineering decision to separate them

Other ray tracing features

- Fresnel reflection and refraction
- Area light sources / soft shadows
- Super sampling for anti aliasing & jittering
- Depth of field blur / real lens modelling
- Motion blur
- Transparency and compositing
- Cartoon shaders
- Texture maps (colour, displacement, normal)
- Other implicit surfaces (metaballs)
- Subdivision surfaces / parametric surfaces
- Constructive solid geometry
- Perlin Noise

Constructive Solid Geometry

- Exactly like intersecting a cube using 3 slabs
- Assemble binary set operations in a tree (works nicely with a dag scene graph)
- Ray intersection class must maintain a list of *all* intersections instead of closest intersection.



Meta Balls

- Implicit function defines surface

$$\sum_{i=0}^n \text{metaball}_i(x, y, z) \leq \text{threshold}$$

$$f(x, y, z) = 1 / ((x - x_0)^2 + (y - y_0)^2 + (z - z_0)^2)$$


- Other better choices for functions too!

- *How do you find the intersection?*
- *What will be the normal?*





Perlin Noise

- Useful for clouds, marble, and other random varying effects
- Code is very small... noise is easy to compute!





Super Sampling and Jittering Samples

- Cast more rays through sub pixel positions.
- Watch out for aliasing!
 - Jittering samples will replace aliasing with noise, which is preferable
 - Noise added to sub-pixel position should not be too large!
- Super sampling can be done adaptively based on scene complexity at the pixel



Depth of Field, Area Lights

- Cast more rays!
 - Sample different random positions on lens aperture by changing the eye position
 - For an area light source, do shadow ray test and lighting computation for a different random positions on the area light source.




Motion Blur

- Knowing object trajectories over a time interval, cast rays at different times and average.
- Can also just average a sequence of images, but this would produce temporal aliasing!
- Can also be smart about not casting few rays where nothing is in motion.



Reflection and Refraction






n_1 = index of refraction of original medium

n_2 = index of refraction of new medium

$$n = \frac{n_1}{n_2}$$



$$c_2 = \sqrt{1 - n^2 \times (1 - c_1^2)}$$

$$R_r = (n \times V) + (n \times c_1 - c_2) \times SN$$



Caustics

- Reflection and refraction scatters light too



Anisotropic BRDFs

- Can use a varying reflectance function for computing the amount of reflected light.



Transparency and Compositing

- Render an image and then combine it with a real photograph.
- Difficult to get lighting correct
 - Typically take a “light probe”, i.e., a photo of a shiny sphere
 - Use sphere image like an area light, and try to address the expensive sampling problem through importance sampling




Cartoon Shaders

- Change lighting computations so that they produce a smaller number of colours.
- Add black edges at boundaries between objects.



Texture Maps

- Can download texture images from a variety of repositories.
- Easy to texture certain primitives (cubes, spheres, planes)
- Meshes often have texture coordinates attached
- ***Sampling is an issue!!***
 - Resolution may be too high or too low giving a poor quality result



Review and More Information

- Textbook
 - 12.3 Spatial Data Structures
 - 13.1 Transparency and Refraction
 - 13.2 Instancing
 - 13.3 Constructive Solid Geometry
 - 13.4 Distributed Ray Tracing
 - Antialiasing, soft shadows / area lights, depth of field, glossy reflection (motion blur not in text but similar)